RAPIDS

The Platform Inside and Out
Release 0.18

@RAPIDSai

https://github.com/rapidsai

https://rapids-goai.slack.com/join

RAPIDS

https://rapids.ai

Open Source Data Science Ecosystem

Familiar Python APIs

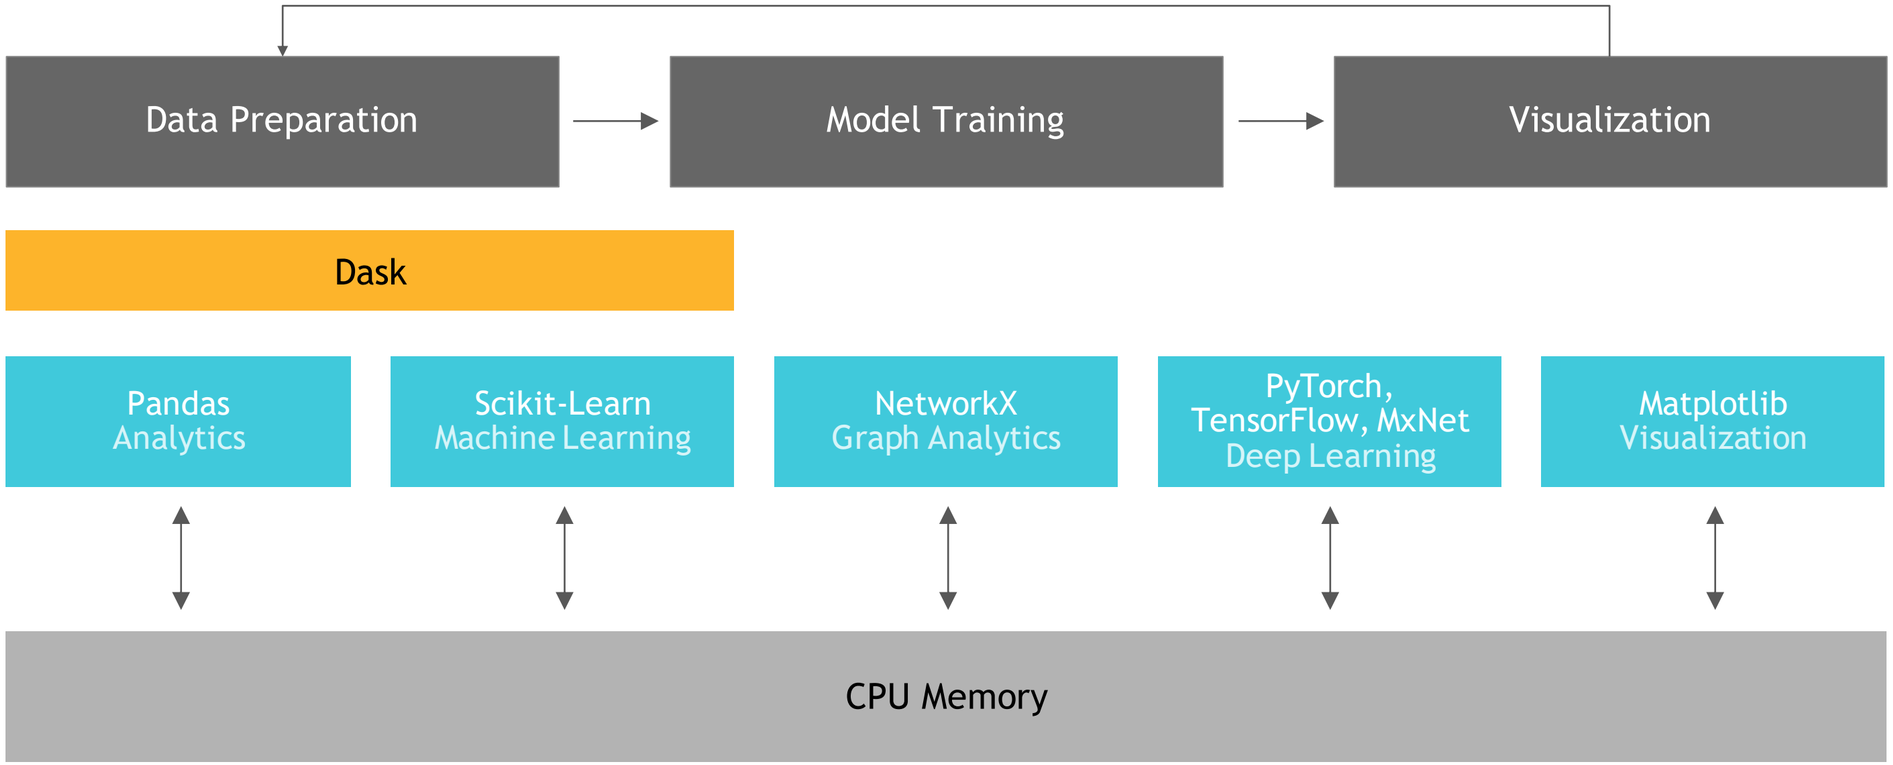

RAPIDS

End-to-End GPU Accelerated Data Science

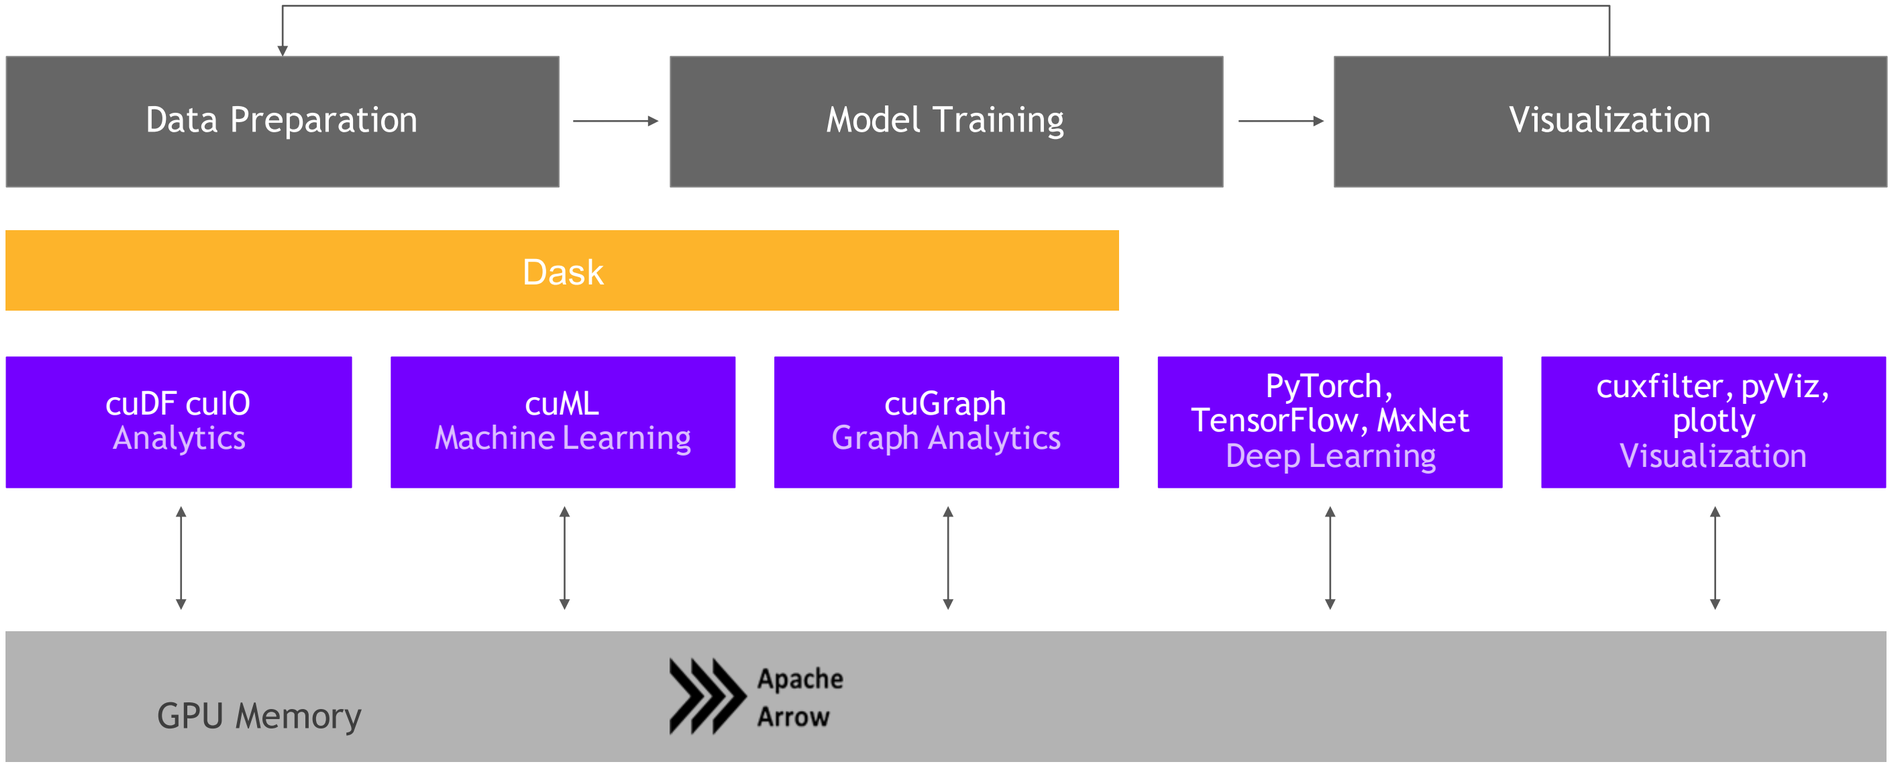

Data Processing Evolution

Faster Data Access, Less Data Movement

Hadoop Processing, Reading from Disk

25-100x Improvement
Less Code
Language Flexible
Primarily In-Memory

Spark In-Memory Processing

5-10x Improvement
More Code
Language Rigid
Substantially on GPU

Traditional GPU Processing

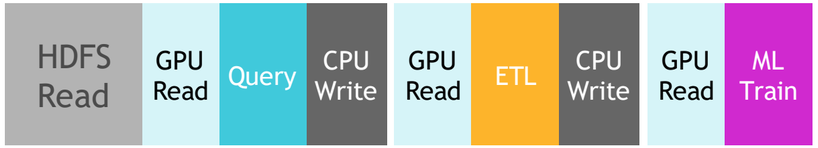

Data Processing Evolution

Faster Data Access, Less Data Movement

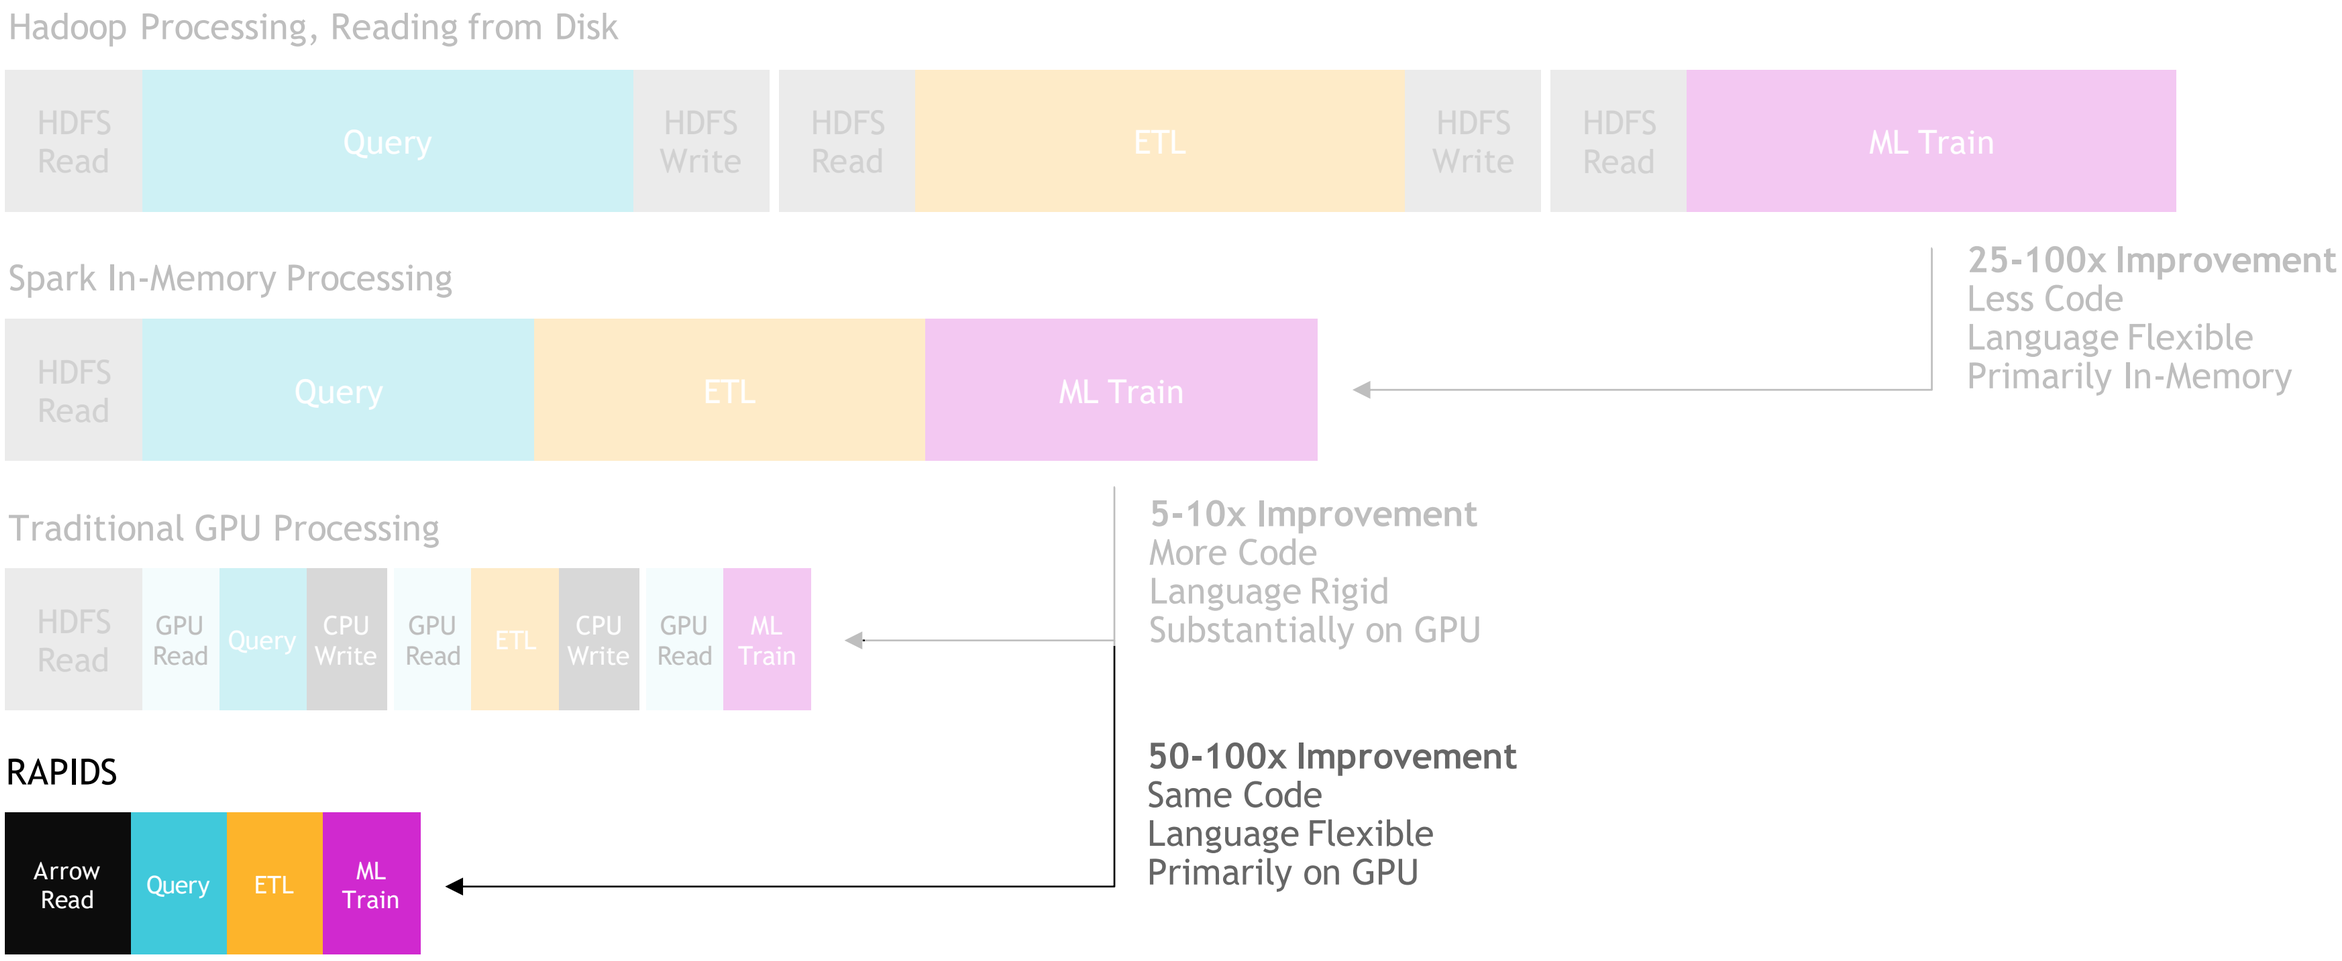

Faster Speeds, Real World Benefits

Faster Data Access, Less Data Movement

cuIO/cuDF - Load and Data Preparation

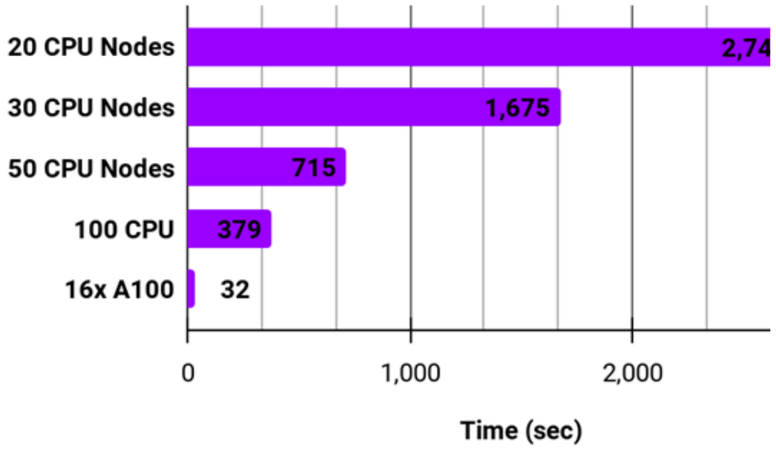

XGBoost Machine Learning

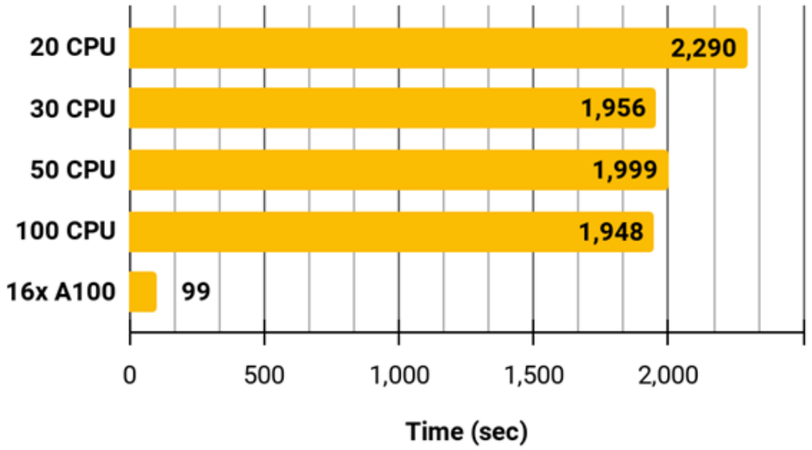

End-to-End

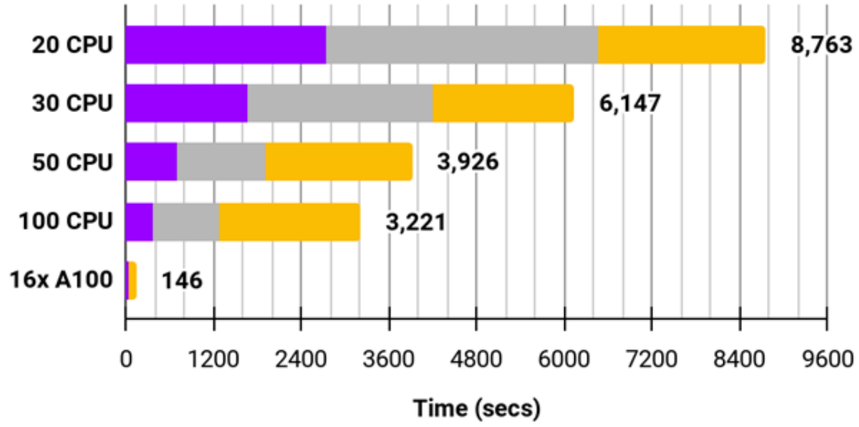

Time in seconds (shorter is better)

cuIO/cuDF (Load and Data Prep) | Data Conversion | XGBoost

Benchmark

200GB CSV dataset; Data prep includes joins, variable transformations

CPU Cluster Configuration

CPU nodes (61 GiB memory, 8 vCPUs, 64-bit platform), Apache Spark

A100 Cluster Configuration

16 A100 GPUs (40GB each)

RAPIDS Version

RAPIDS 0.17

Scale Up with RAPIDS

Scale Up / Accelerate

RAPIDS AND OTHERS

Accelerated on single GPU

NumPy -> CuPy/PyTorch/..
Pandas -> cuDF
Scikit-Learn -> cuML
NetworkX -> cuGraph
Numba -> Numba

RAPIDS

PYDATA

NumPy, Pandas, Scikit-Learn, NetworkX, Numba and many more

Single CPU core
In-memory data

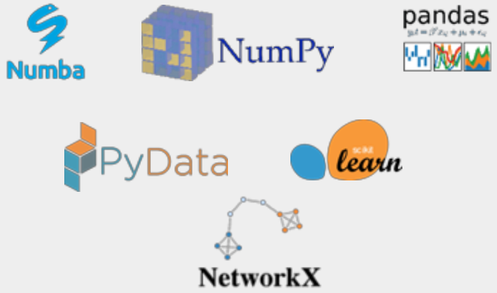

Scale Out with RAPIDS + Dask with OpenUCX

Scale Up / Accelerate

RAPIDS AND OTHERS

Accelerated on single GPU

NumPy -> CuPy/PyTorch/..
Pandas -> cuDF
Scikit-Learn -> cuML
NetworkX -> cuGraph
Numba -> Numba

RAPIDS

RAPIDS + DASK WITH OPENUCX

Multi-GPU
On single Node (DGX)
Or across a cluster

RAPIDS

DASK

PYDATA

NumPy, Pandas, Scikit-Learn, Numba and many more

Single CPU core
In-memory data

Numba NumPy pandas PyData learn NetworkX

DASK

Multi-core and distributed PyData

NumPy -> Dask Array
Pandas -> Dask DataFrame
Scikit-Learn -> Dask-ML
… -> Dask Futures

DASK

Scale Out / Parallelize

cuDF

RAPIDS

GPU Accelerated Data Wrangling and Feature Engineering

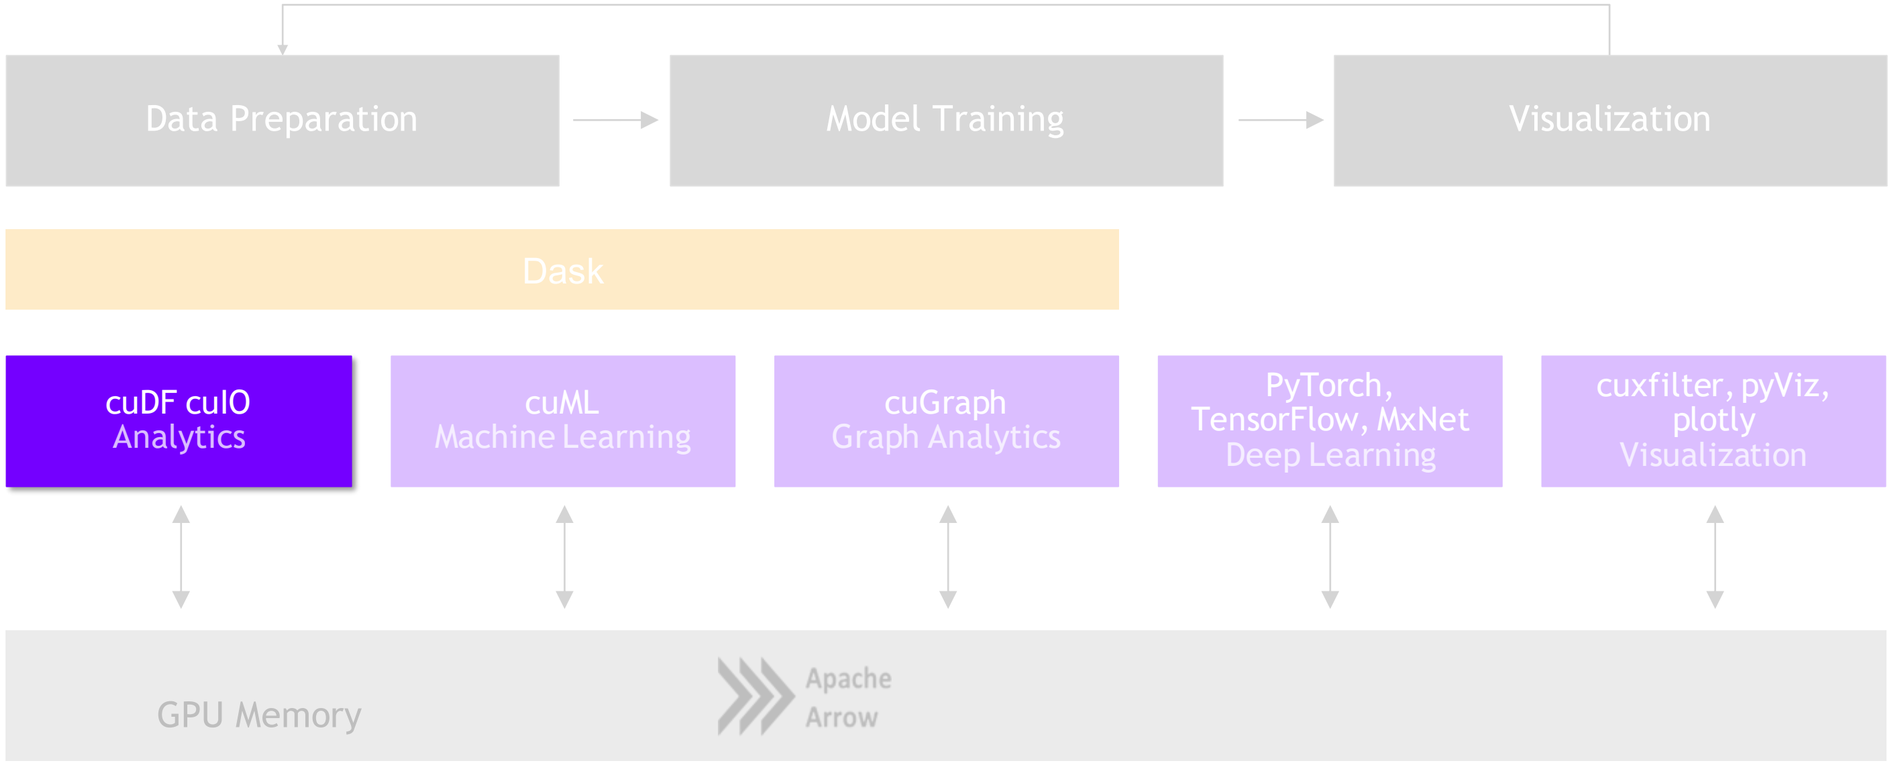

ETL - the Backbone of Data Science

cuDF is…

PYTHON LIBRARY

- A Python library for manipulating GPU DataFrames following the Pandas API
- Python interface to CUDA C++ library with additional functionality
- Creating GPU DataFrames from Numpy arrays, Pandas DataFrames, and PyArrow Tables
- JIT compilation of User-Defined Functions (UDFs) using Numba

```
In [2]: #Read in the data. Notice how it decompresses as it reads the data into memory.
gdf = cudf.read_csv('/rapids/Data/black-friday.zip')
```

```
In [3]: #Taking a look at the data. We use "to_pandas()" to get the pretty printing.
gdf.head().to_pandas()
```

Out[3]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Ca
0	1000001	P00069042	F	0-17	10	A	2	0	3
1	1000001	P00248942	F	0-17	10	A	2	0	1
2	1000001	P00087842	F	0-17	10	A	2	0	12
3	1000001	P00085442	F	0-17	10	A	2	0	12
4	1000002	P00285442	M	55+	16	C	4+	0	8

```
In [6]: #grabbing the first character of the years in city string to get rid of plus sign, and converting
 to int
gdf['city_years'] = gdf.Stay_In_Current_City_Years.str.get(0).stoi()
```

```
In [7]: #Here we can see how we can control what the value of our dummies with the replace method and turn
strings to ints
gdf['City_Category'] = gdf.City_Category.str.replace('A', '1')
gdf['City_Category'] = gdf.City_Category.str.replace('B', '2')
gdf['City_Category'] = gdf.City_Category.str.replace('C', '3')
gdf['City_Category'] = gdf['City_Category'].str.stoi()
```

Benchmarks: Single-GPU Speedup vs. Pandas

cuDF v0.13, Pandas 0.25.3

- Running on NVIDIA DGX-1:
  - GPU: NVIDIA Tesla V100 32GB
  - CPU: Intel(R) Xeon(R) CPU E5-2698 v4 @ 2.20GHz
- Benchmark Setup:
  - RMM Pool Allocator Enabled
  - DataFrames: 2x int32 columns key columns, 3x int32 value columns
  - Merge: inner; GroupBy: count, sum, min, max calculated for each value column

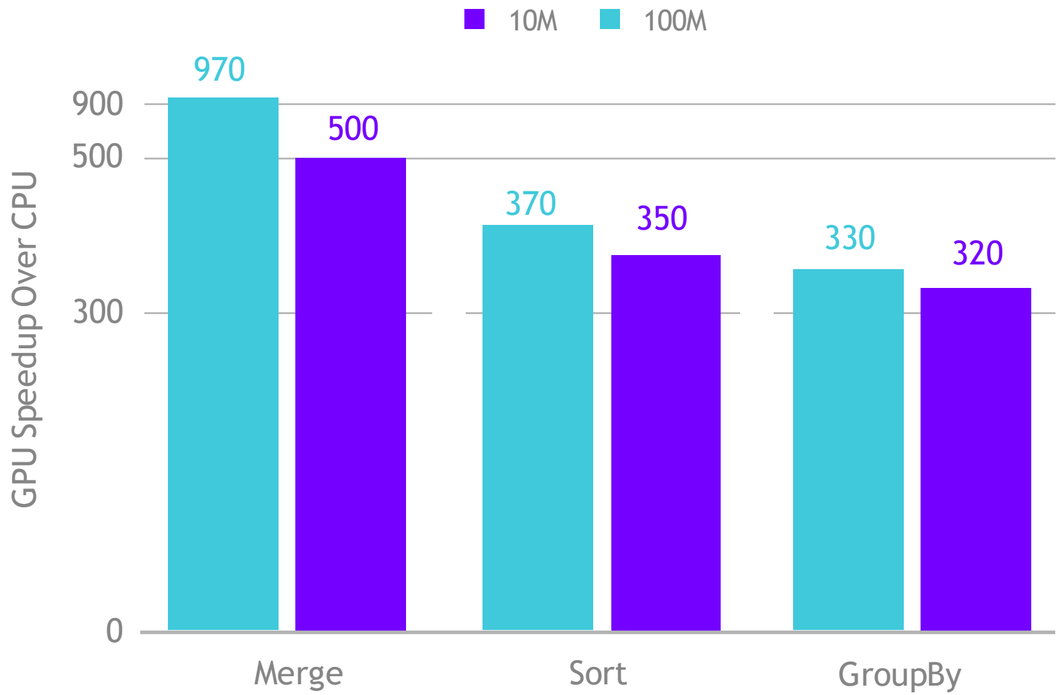

cuML

Machine Learning

More Models More Problems

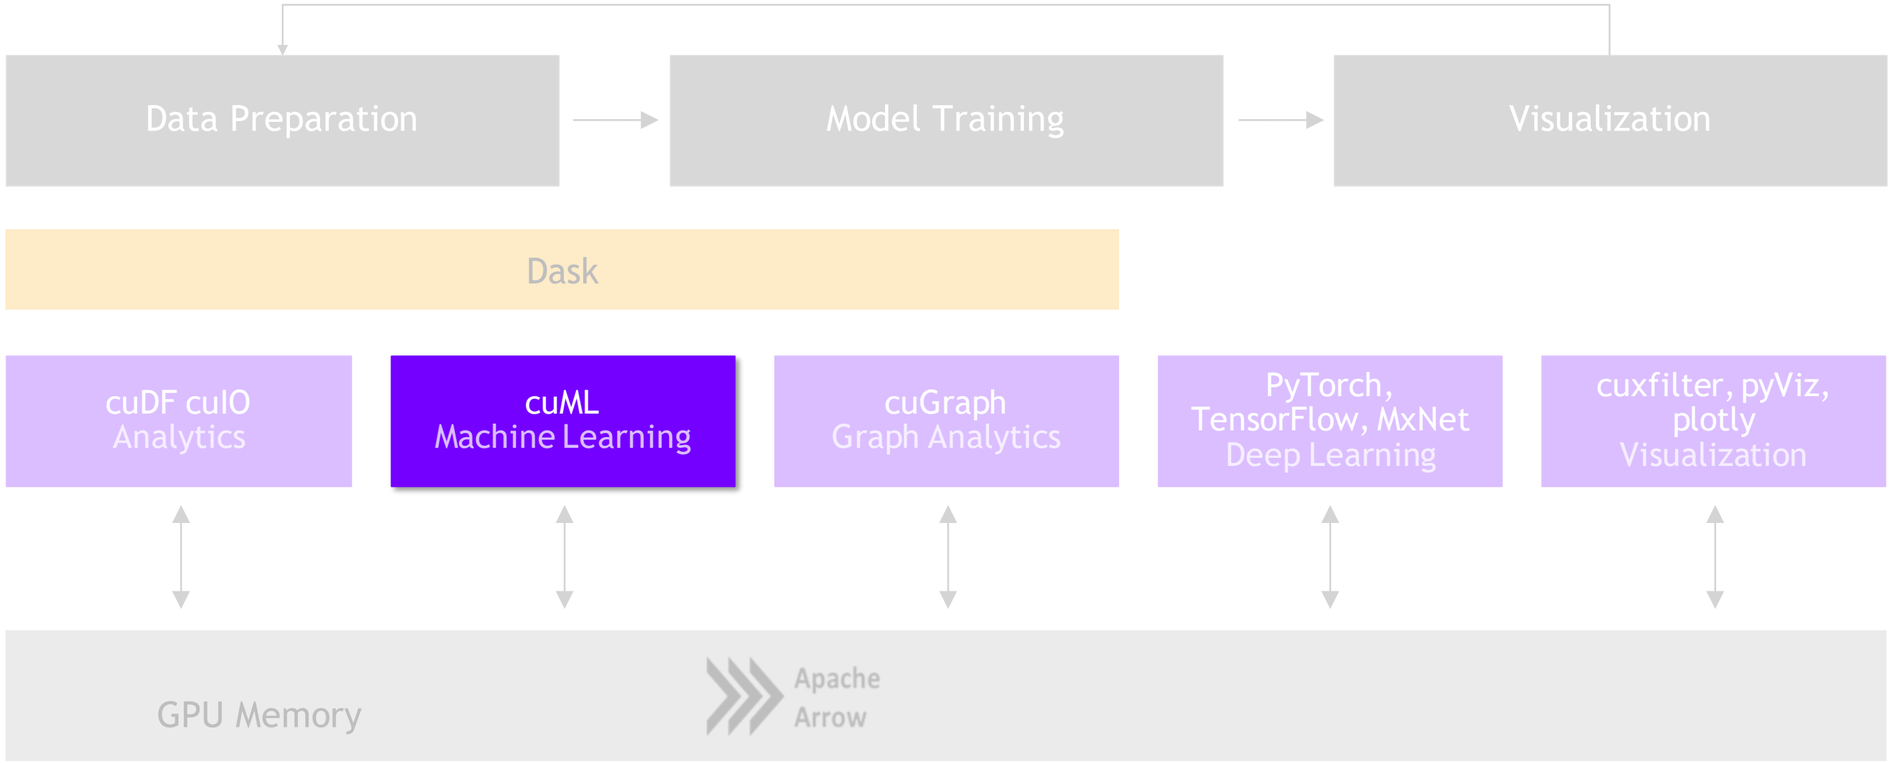

RAPIDS Matches Common Python APIs

CPU-based Clustering

```
from sklearn.datasets import make_moons
import pandas


X, y = make_moons(n_samples=int(1e2),
                  noise=0.05, random_state=0)

X = pandas.DataFrame({'fea%d'%i: X[:, i]
              for i in range(X.shape[1])})
```

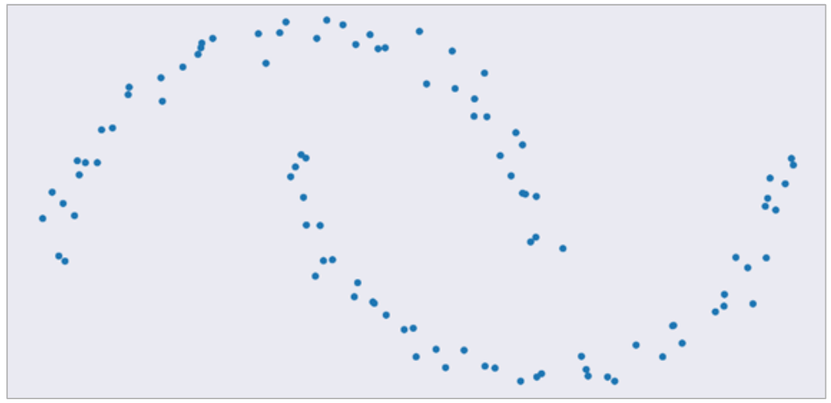

```
from sklearn.cluster import DBSCAN
dbscan = DBSCAN(eps = 0.3, min_samples = 5)


y_hat = dbscan.fit_predict(X)
```

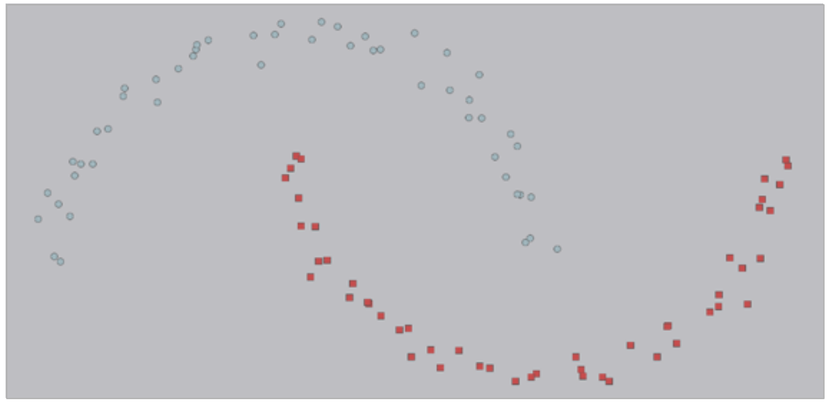

RAPIDS Matches Common Python APIs

GPU-accelerated Clustering

```
from sklearn.datasets import make_moons
import cudf


X, y = make_moons(n_samples=int(1e2),
                  noise=0.05, random_state=0)

X = cudf.DataFrame({'fea%d'%i: X[:, i]
                    for i in range(X.shape[1])})
```

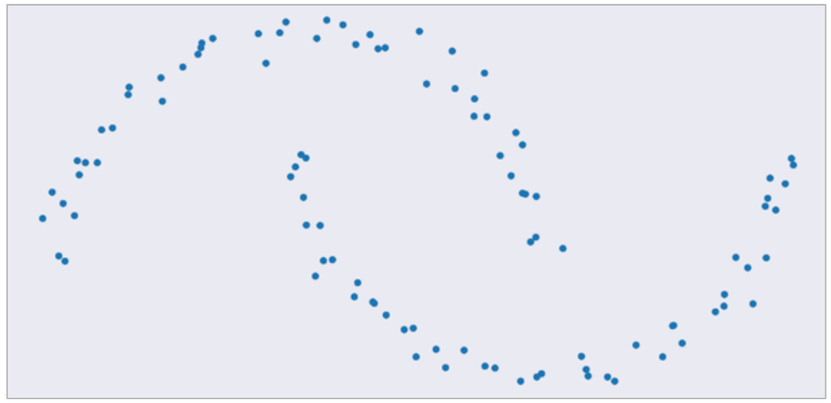

```
from cuml import DBSCAN
dbscan = DBSCAN(eps = 0.3, min_samples = 5)


y_hat = dbscan.fit_predict(X)
```

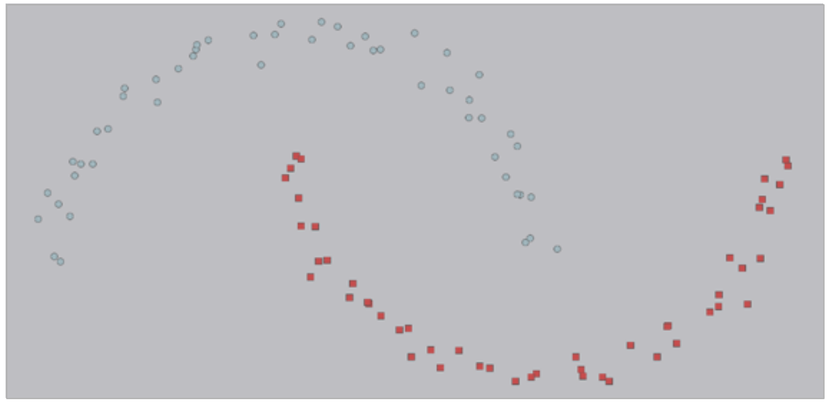

Algorithms

GPU-accelerated Scikit-Learn

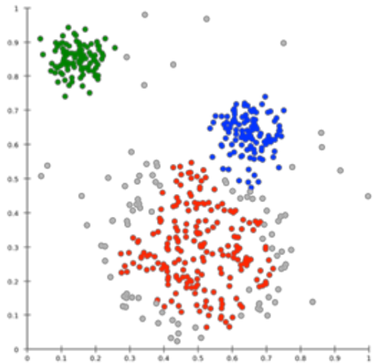

Cross Validation

Hyper-parameter Tuning

More to come!

Classification / Regression	Decision Trees / **Random Forests** Linear/Lasso/Ridge/**LARS**/ElasticNet Regression Logistic Regression K-Nearest Neighbors (**exact or approximate**) Support Vector Machine Classification and Regression Naive Bayes
Inference	Random Forest / GBDT Inference (FIL)
Preprocessing	Text vectorization (TF-IDF / Count) Target Encoding Cross-validation / splitting
Clustering Decomposition & Dimensionality Reduction	K-Means **DBSCAN** Spectral Clustering Principal Components (including iPCA) Singular Value Decomposition UMAP Spectral Embedding T-SNE
Time Series	Holt-Winters Seasonal ARIMA / Auto ARIMA

Benchmarks: Single-GPU cuML vs Scikit-learn

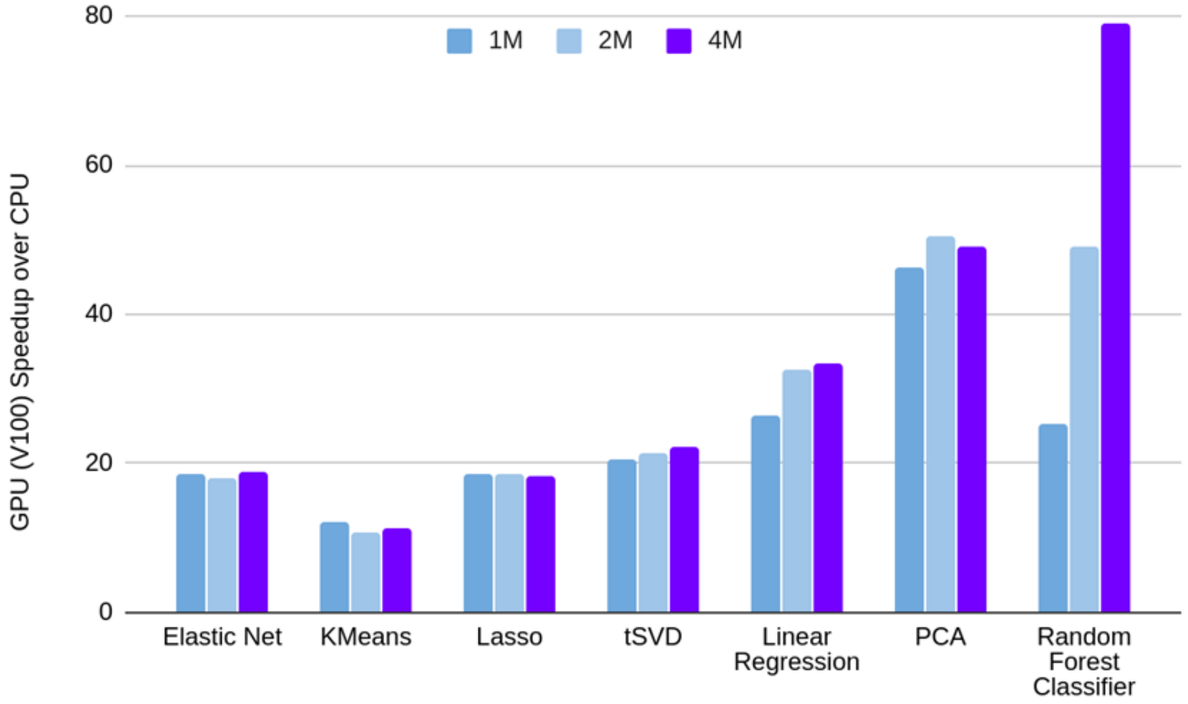

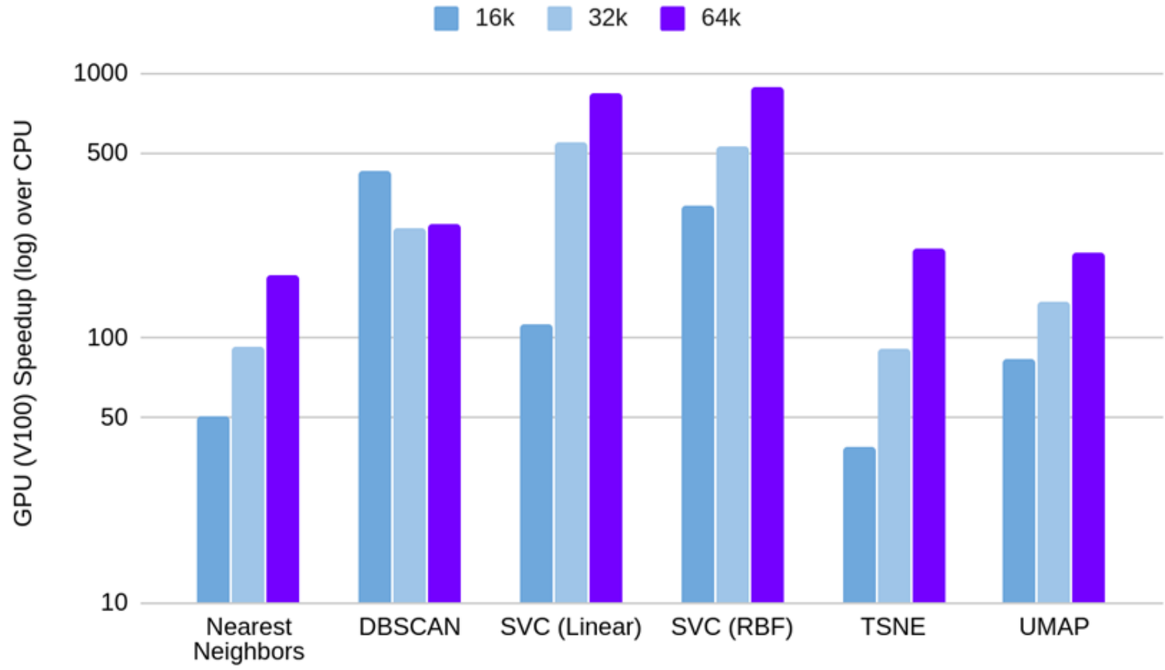

1x V100 vs. 2x 20 Core CPUs (DGX-1, RAPIDS 0.15)

Getting Started

Easy Installation

Interactive Installation Guide

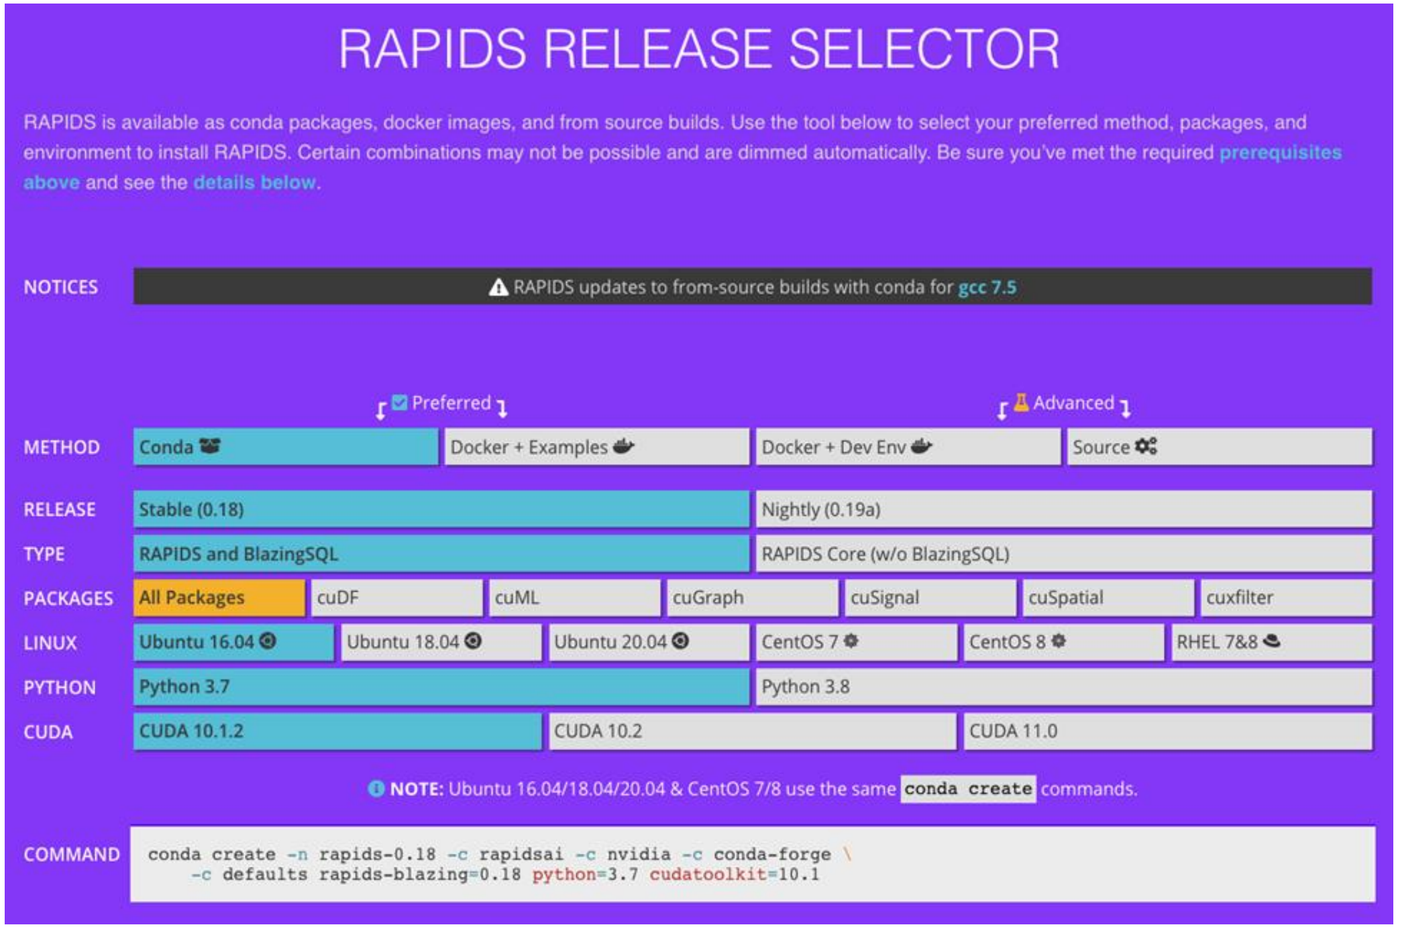

https://rapids.ai/start.html

RAPIDS Docs

Up to Date and Easy to Use

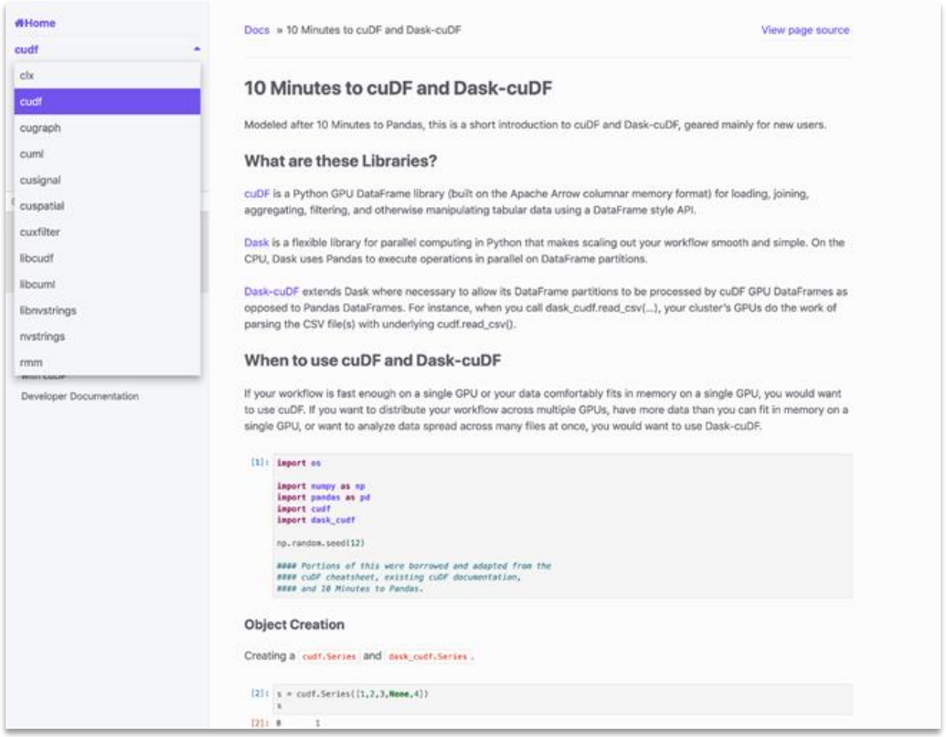

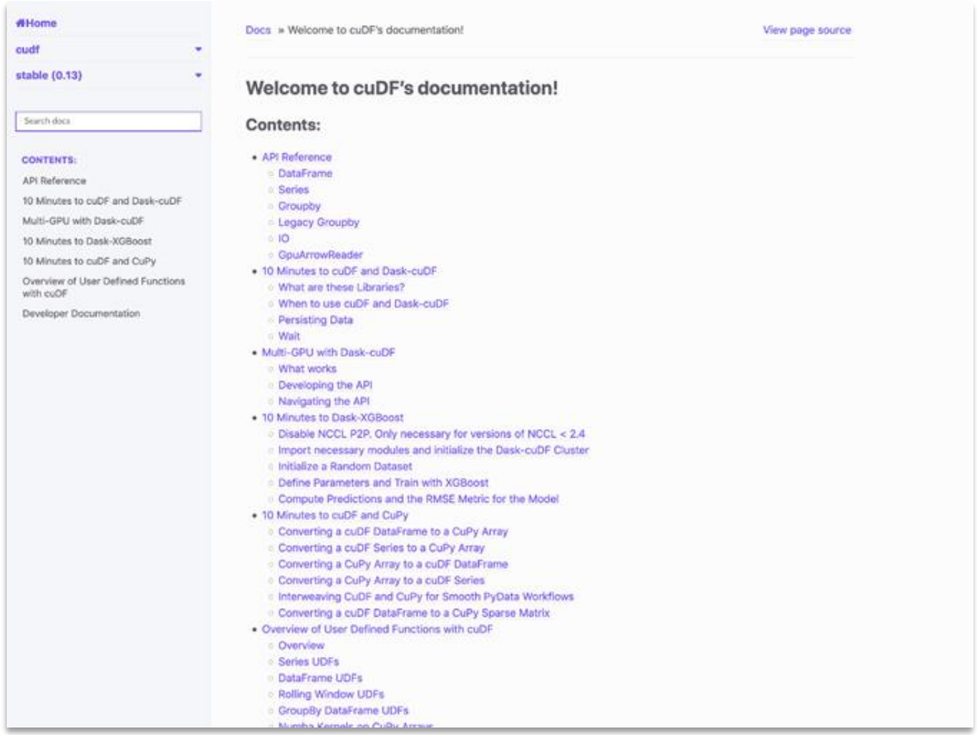

https://docs.rapids.ai